

## ASSIGNMENT-2

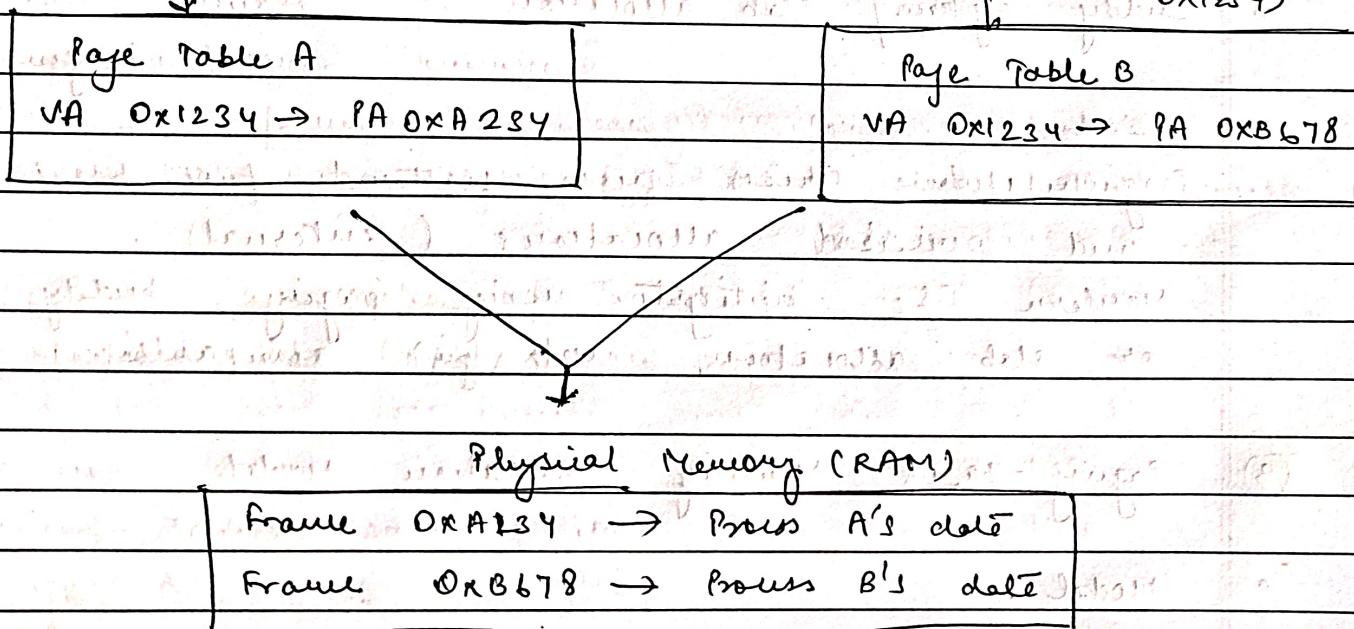
(1)

In Modern systems, each process uses its own logical (virtual) addresses. These are translated into physical addresses in main memory through the Memory Management Unit (MMU) using a page table.

- CPU generates a logical address (virtual address).
- MMU consults the page table to find the corresponding physical frame.
- Physical address is then accessed in RAM.
- This allows multiple processes to coexist safely without interfering with each other.

Illustration

Process A (Virtual Address: 0x1234)      Process B (Virtual Address: 0x1234)



Modern systems use virtual memory where each process generates a logical address. The MMU translates it.

into a physical address using page tables, ensuring isolation between processes.

## (2) Memory layout

P1 (10)	Free(5)	P2 (20)	Free(10)	P3 (15)
---------	---------	---------	----------	---------

- External fragmentation: 15 MB free but scattered → can't fit a 12 MB process.
- Internal fragmentation: If P4 needs 18 MB, but gets 20 MB, 2 MB wasted inside.

## Mitigation (beyond compaction).

- Modern OS uses:
- Paging → avoids external fragmentation (only last page is wasted).
- Buddy system / slab allocation → reduces internal fragmentation.

Fragmentation arises from scattered free blocks (external) and oversized allocations (internal).

Modern OS mitigate using paging, buddy system, or slab allocation, not just compaction.

## (3) Paging-based memory Allocation Model

- Model:
  - Divide physical memory into fixed-size frames (e.g. 4 kB).
  - Divide process memory into same-sized pages.
  - use a page table to map virtual pages → physical frames.

- o Trade-off:

- No external fragmentation (pages can go anywhere)
- Internal fragmentation possible in the last page.
- Memory overhead (page tables consume extra space)
- Speed penalty (extra translating step, initialized by TLB cache).

### (4) OS - Hardware Interaction in Virtual Memory

- o Process:

- CPU generates a virtual address
- MMU (Memory Management Unit) translates it to a physical address via page tables.
- TLB (Translation Lookaside Buffer) caches recent translations → speed up access
- Page Fault: If page not in RAM, OS handles by fetching from disk (swap).

- o Key Hardware Structure:

- MMU → performs translation & enforces protection
- Pages Table Base Register (PTBR) → points to current process page table.
- TLB → fast cache of translations.
- Protection bits (R/W/X) in page table

### (5)

Given:

- Virtual address size = 16 bits

- Page size = 1KB =  $2^{10}$  bytes

- Page Table entry (PTE) size = 2 bytes.

a.

Number of Virtual PagesVirtual address space  $= 2^{16} = 65,536$  bytes (64 KB).Page size  $= 2^{10} = 1024$  bytes.Number of virtual pages  $= \frac{\text{Virtual Address Space}}{\text{Page size}} = \frac{2^{16}}{2^{10}} = 64$  $\therefore 64$  virtual pages.

b.

Pages Table Size

Each virtual page needs 1 entry

Total entries  $= 64$ Size  $= 64 \times 2$  bytes  $= 128$  bytes $\therefore$  Pages table size  $= 128$  bytes

(6)

Allocation (all 3 methods)

P1  $\rightarrow$  212 KB, P2  $\rightarrow$  417, P3  $\rightarrow$  1121 KB allocated.P4  $\rightarrow$  cannot fit (largest free block  $= 259$  KB  $< 426$  KB)Used unused memory:First-fit  $= 259$  KBBest-fit  $= 259$  KBWorst-fit  $= 259$  KB

All three give same utilization here (741 KB used, 259 KB free) since processes arrived in order and split the single block idly.

(7)

Page-replacement (reference freq = 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2 frames).

a & b → simulation result (summary):

FIFO : page faults = 10

Optimal : page faults = 7

LRU : page faults = 9.

(Short justification: Optimal by definition produces the minimum fault (7). FIFO replaces the oldest page and here causes 10 faults; LRU (replaces least recently used) gives 9 faults).

c → comment (Belady's Anomaly):

- Optimal and LRU are slack algorithms (they do not suffer Belady's anomaly) and tend to perform better.
- FIFO is not a slack algorithm and can exhibit Belady's anomaly (faults can increase when frames decrease). In this case FIFO performs worse than Optimal and LRU, illustrating that naive replacement (FIFO) may give suboptimal behaviour - exactly the kind of effect Belady described.

(8)

Dirty - page overhead and optimization

Given: disk write = 10ms per page, memory write = 100ms. 30% of replaced pages are dirty, replacing 1000 pages.

- a — extra time due to dirty pages;
- Number dirty =  $0.80 \times 1000 = 800$  pages.
  - Time for dirty writes =  $800 \times 10 \text{ ms} = 8000 \text{ ms} = 8 \text{ seconds}$   
(for  $10^3 \text{ s} = 3,000,000 \text{ ms}$ , which equals  $30,000,000$  memory writes at 100 ms each → shows disk I/O dominates).

- b — Optimization to reduce overhead;

- use asynchronous write-behind (page cleaner / write-back mechanism)
- A background page-cleaner writes dirty pages to disk ahead of eviction so the page fault handler rarely blocks on disk I/O.
- combine with an eviction policy that prefers clean pages (e.g. enhanced CLOCK / Second-chance that tracks dirty bits).
- Result: fewer synchronous disk writes during replacement → much lower page fault latency.

### (g) a) Keeping mission-critical task from thrashing.

- Working-set tracking: measures each task's working set. The scheduler / VM subsystem pins or reserves the estimated working set for critical tasks so they keep their pages in RAM and don't suffer page faults.

- Priority - aware replacement: use a working-set enhanced - clock replacement that combines working-set age checks with processing priority: when reclaiming, prefer to evict pages belonging to low-priority processes (background) and only reclaim a virtual task's pages if its working set has genuinely cooled. Also prefer dirty clean pages and background-clean dirty pages to avoid thrashing.
- Admission & prefetching: refuse to admit extra high-memory best-effort tasks when virtual task's working set cannot be satisfied; prefetch pages for virtual task (e.g., model weights) at startup to avoid runtime faults.

### b) Recommended memory allocation strategy & justification

- Hybrid reservation + dynamic pool:
  1. Reserved real-time pool: statically / periodically reserve a fraction of RAM for mission-critical processes (unlock/rtcgroup + preallocated pools or huge pages).
  2. Shared dynamic pool: remaining memory is managed by the normal VM with priority-aware reclaim for best-effort tasks.
- Supporting techniques: use huge pages for large model data (reduce TLB pressure), slab/pool allocators for predictable kernel allocations, an active page-cleaner daemon, and cgroup.

memory QoS to throttle noncritical workloads.

- Why this balances both the reserved pool guarantees real-time responsiveness and predictable latency for safety-critical modules while the dynamic pool maximizes overall utilization by allowing noncritical services to expand / contract and absorb reeling / so cost - combining priority-aware replacement and prefetching minimizes blocking I/O and keeps system efficient.